

Using Generative Neural Networks to generate artist-inspired artwork

Mihai Bojescu

*Master in Artificial Intelligence and optimisation
Faculty of Computer Science
University “Alexandru Ioan Cuza” of Iași
Iași, Romania
bojescu.mihai@gmail.com*

Radu Șolcă

*Master in Artificial Intelligence and optimisation
Faculty of Computer Science
University “Alexandru Ioan Cuza” of Iași
Iași, Romania
radu.ssolca@gmail.com*

Abstract—This document contains a study on how Generative Neural Networks could be applied in order to transform noise into artwork inspired by famous artists. The architecture of the Network consists of a Convolutional Neural Network for the discriminator and a U-Net Convolutional Neural Network for the generator. The data on which the network is trained is a dataset with famous paintings by Jean Monnet.

Index Terms—Generative Neural Networks, GANs, Convolutional Neural Networks, U-Nets.

I. INTRODUCTION

In this document, we will study how GANs could be used to generate pieces of art in the Jean Monnet style. GANs have shown remarkable success in generating realistic images by using random noise as input, and we aim to use them as an application for producing fine art.

Monnet’s paintings are characterised by their vibrant colors and expressive brushwork, and we will try to replicate this signature in the pictures that the GAN produces, effectively creating a ‘digital Monnet’. The model will play by rotation the role of an art critic and an artist in order to improve itself until - as an art critic - it cannot distinguish between real, pre-existing art and newly created art.

II. GANs

Generative Adversarial Networks, a type of Neural Network that was originally built in 2014 by Ian Goodfellow and his colleagues in June 2014 [1] is a type of machine learning system that involves two neural networks competing against each other. The competition can be seen as a game where two teams have to play a zero-sum game: one team’s win is the other team’s loss. The system is part of the unsupervised learning category, but can also be used for semi-supervised learning, fully-supervised learning and for reinforcement learning. The aforementioned system can be used for a vast set of use-cases, ranging from image generation - the classical use-case - to music generation.

GANs are known to produce high-quality results with good performance, but suffer of low diversity.

In our paper, we present a classical use-case for GANs: image generation. As seed, we will be using random noise from the uniform distribution and we will use the seed to generate art in the Monet style.




Fig. 1: Data generation problem

A. Data processing

In our project, we have used only a resize transformation, which resizes the images from 256x256 to 32x32.

B. Discriminator

For a discriminator, we will use a classical cone-shaped convolutional neural network that receives an image with 3 channels and in turn creates, 64, 128, 256, 512 feature maps, which are later flattened to a 1-dimensional array. The network makes use of a Leaky Rectified Linear Unit activation function for the feature maps building, and a Sigmoid activation function at its output.

The trainer of the discriminator uses - for each time it is used - a batch of real images and a batch of fake ones, which are used to “teach” the model how to classify images: fake ones are considered negative, while real ones as positive. For training we use the Wasserstein distance - also named Earth Mover’s distance - to compute the loss of the discriminator [2]:

$$\begin{aligned} loss_D = & - \frac{\sum_{i=0}^n discriminator_output_{real_image}^i}{n} \\ & + \frac{\sum_{i=0}^n discriminator_output_{fake_image}^i}{n} \\ & + gradient_loss(fake_image, real_image) \quad (1) \end{aligned}$$

Where *gradient_loss* is a function that computes, given a real and a fake image, a penalty that encourages the gradients

of the critic to have norm 1 almost everywhere. This is used in order to enforce the Lipschitz continuity constraint, which in turn should improve the results.

C. Generator

The generator of the GAN is comprised of a classical U-Net architected CNN. U-Nets are well known for producing pictures with the same shape as the input, which in our case proved beneficial, as we could provide noise for each pixel of the image, noise. After training, the generator *should* output realistic-looking Monet-inspired pictures.

For its architecture, we used multiple layers of convolutional and max-pooling layers that downscale the image and extract features, data which is later used to upscale the image back using multiple transposed convolutional and classical convolutional layers. For the activation function we also use the Leaky Rectified Linear Unit function.

The trainer performs the following steps for the generator:

- 1) Using a given noise tensor, run the generator
- 2) Using the output of the generator, run the discriminator
- 3) Use the inverted results of the discriminator to update the loss of the generator

$$loss_G = -\frac{\sum_{i=0}^n discriminator_output_{fake_image}^i}{n} \quad (2)$$

D. Tuning

In order to tune our model, we have used Weights & Biases and have performed multiple sweeps. We have used these sweeps to tune the parameters in order to reach the results specified in the Results (IV) section. Most of our runs could show that the losses of the generator and discriminator converge quite fast to 0.

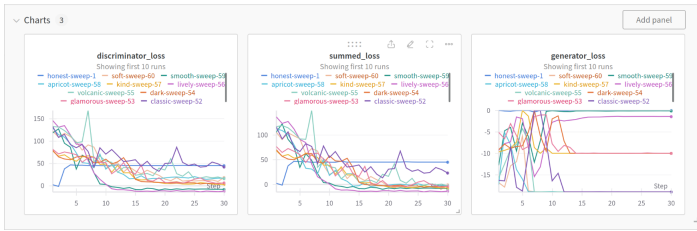


Fig. 2: Weights & Biases run

III. SHORTFALLS

GANs are well-known for being unstable in training, and are also well-known for suffering of the mode collapse problem - learning so well the training data of the discriminator, that it becomes a generator that outputs only the training data. In our case, the diversity issue was mitigated by using the gradient penalty function. The model proved to be highly unstable, most often outputting noise.

IV. RESULTS

In general, most of the results could not be distinguished from the random noise used to generate the generator's outputs (3, 4). The red in the discriminator output means that the discriminator classified the output of the generator as a fake image. There were results where some shapes could be distinguished, but they were not close to the original paintings (5, 6).

We trained our model on an Intel Core i7 1260p processor, a NVIDIA RTX 3070Ti laptop GPU and a NVIDIA T4 GPU in the cloud. We collected the following metrics from our training runs:

Processor	Mean timing per iteration
Intel i7 1260p	12s
NVIDIA RTX 3070Ti	1.3s
NVIDIA T4	1.83s

TABLE I: Performance metrics

V. CONCLUSIONS

In literature, GANs can be used to generate art from random noise. In our project we aimed to replicate these results, but at best we could generate only shaped artifacts. The model would need further debugging and fine-tuning before it can generate meaningful data.

REFERENCES

- [1] I. J. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, Y. Bengio (2014). "Generative Adversarial Nets". Département d'informatique et de recherche opérationnelle Université de Montréal Montréal, QC H3C 3J7. arXiv: 1406.2661v1
- [2] M. Arjovsky, S. Chintala, L. Bottou (2017). "Wasserstein GAN". Courant Institute of Mathematical Sciences. arXiv: 1701.07875v3




Fig. 3: Noisy sample




Fig. 4: Noisy sample 2




Fig. 5: Sample with shape




Fig. 6: Sample with shape 2